



MSIB32500 Advanced Bioinformatics Spring 2019



RNAseq Data Analysis and Clinical Applications, Part I

Riyue Sunny Bao, Ph.D.

Research Assistant Professor (Bioinformatics)

Center for Research Informatics & Department of Pediatrics

The University of Chicago





Outline

- ▶ Part I (05/25/2019)
 - ▶ Introduction to RNAseq technology and clinical applications
 - ▶ Hands on: From raw data to gene expression quantification
- ▶ Part II (06/01/2019)
 - ▶ Differential gene expression analysis and data visualization
 - ▶ Hands on: Identification of genes and pathways significantly changed under condition
 - ▶ **Homework assignment**
- ▶ Part III (06/08/2019)
 - ▶ How to associate gene expression data with clinical outcome
 - ▶ Hands on: Use gene expression data to discover tumor subtypes and survival analysis



Class materials

- GitHub

- <https://github.com/MScBiomedicalInformatics/MSIB32500>

- This lecture note contains the same contents as the notebook. In addition, the notebook also contains hands-on materials

- **lecture8.pdf**

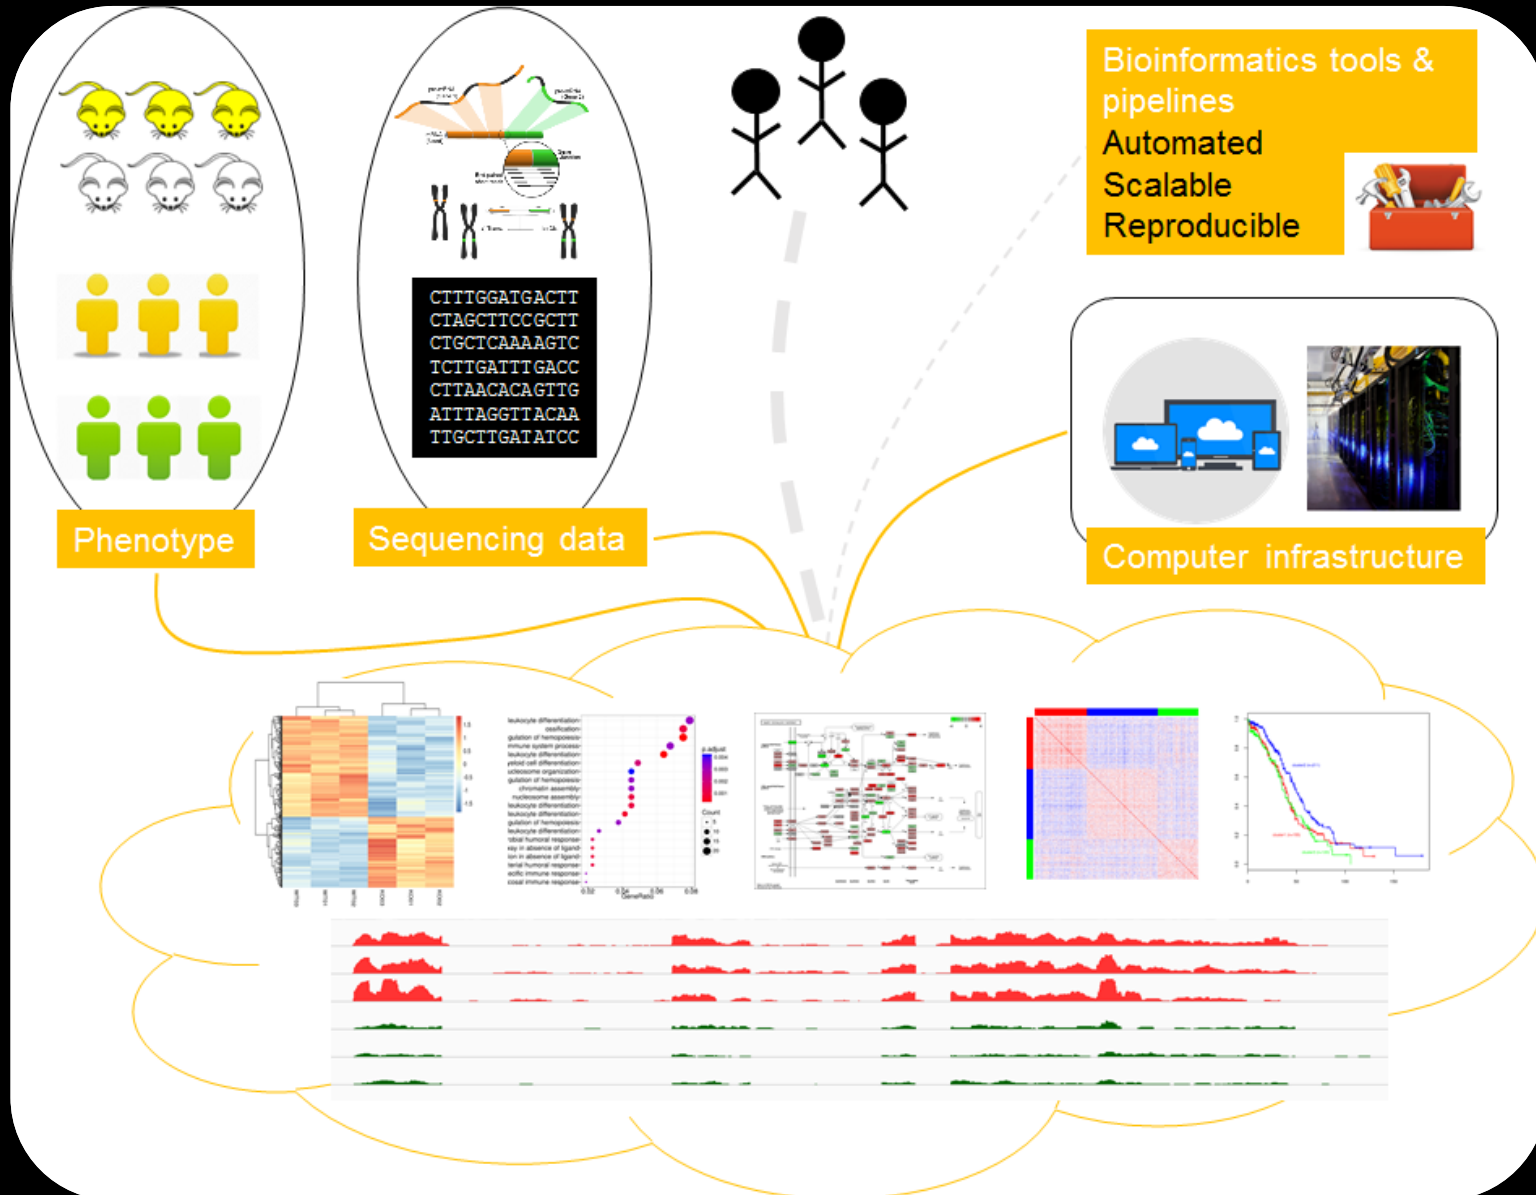
- **Handson8.Rmd**

- Gardner high-performance computing (HPC) clusters (hands on practice)



Objective

- ▶ Learn the good-practice RNAseq analysis pipeline
- ▶ Learn commonly used bioinformatics tools
- ▶ Practice the automated, scalable pipeline
- ▶ Explore the quality metrics and input/output of the RNAseq pipeline
- ▶ Visualize result files and quality plots






Biological and clinical questions

- ▶ I am interested in studying transcriptional landscape shift before and after drug treatment in cell lines
- ▶ I want to identify which pathways are affected after knocking down my favorite gene in mice
- ▶ I have expression data of clinical isolates collected at various time points, when patient's response changed. Why?
- ▶ I have a cohort of patients and want to discover which gene signature predicts patient's response to treatment
- ▶ I want to detect gene fusions, expressed mutations, and disrupted isoforms in tumors that may be related to disease

... and more!

The Sequencing Iceberg

One sequencing
run every 3 days
(per instrument)



28 instruments

Data storage
Data transfer
Data analysis

1 TB of images,
intermediate
analysis files and
tracking data per
run

Current Sequencing Technologies

Illumina

HiSeq 2000/2500/X *TEN*



Illumina

MiSeq



Ion Torrent (Life Technologies)
Ion Proton

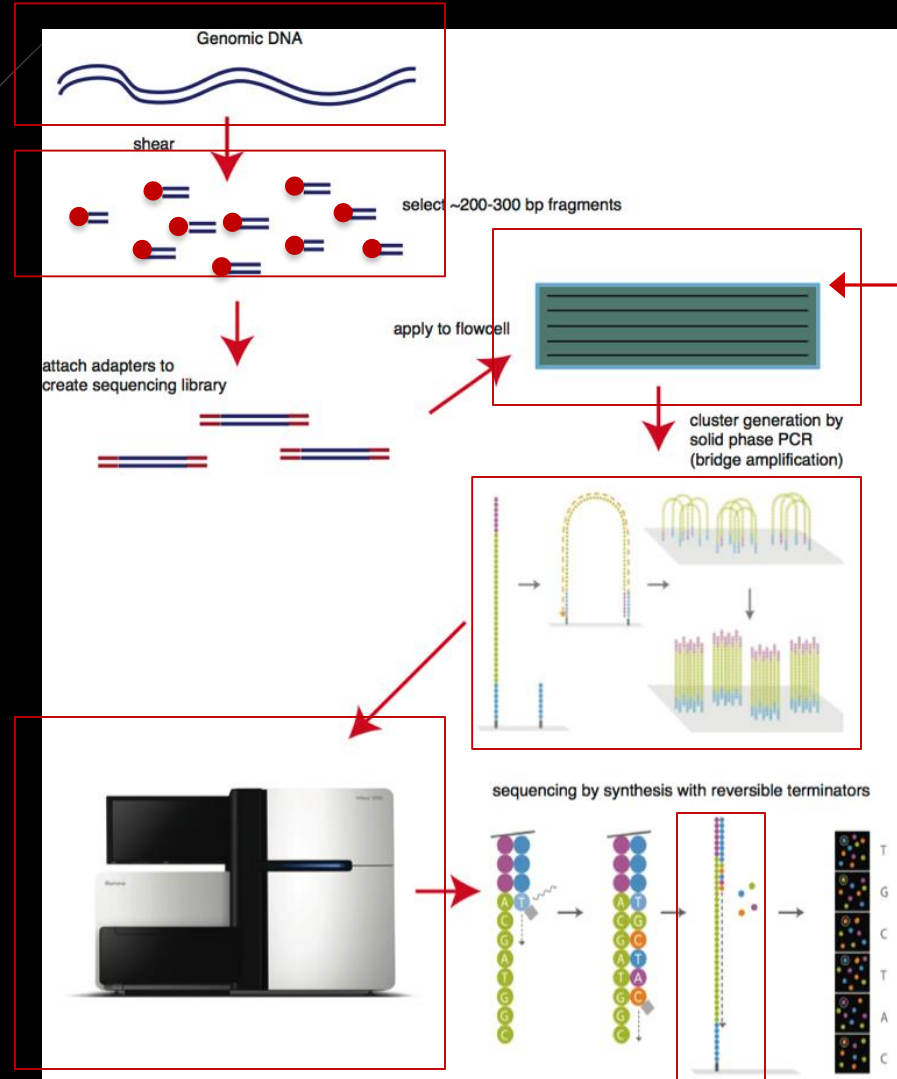


PacBio

RS



Illumina




- Sample
- Library
- DNA fragment
- Barcode
- Run
- Flow cell
- Lane
- Cluster
- Read
- Adapter/Primer

Though we are talking about Illumina here, many of those terms can be applied to other sequencing technologies.

Illumina

DNA fragment


Adapter/Primer



Single-End



Paired-End




- Sample
- Library
- DNA fragment
- Barcode
- Run
- Flow cell
- Lane
- Cluster
- Read
- Adapter/Primer

Fragment size = R1+R2+insert size

What is RNAseq?

High-throughput sequencing of RNA: Profile, identify or assemble transcripts

- Detect gene expression changes between conditions
- Identify novel splice sites / exons, mutations, fusion genes, etc.
- Broad detection range, high sensitivity, low requirement of RNA amount
- Available for all species (reference genome is optional): reference genome-guided alignment or *de novo* assembly



Biological sample or
clinical specimen

Millions of reads!




Which factors to consider if I want to initiate an RNAseq experiment?

- Sample source (fresh cells, frozen, FFPE)
- Sample quality (RNA quality, DNA contamination, tumor purity)
- RNA concentration (e.g. 100ng total RNA)
- Ribosome RNA depletion (accounts for 80% of total RNA)
- **Library type (single-end or paired-end reads)**
- **Library strandness (unstranded or stranded/directional)**
- **Sequencing depth (20 million, 50 million, or >100 million reads)**

Library type: Single-end vs paired-end reads


DNA fragment



Single-End

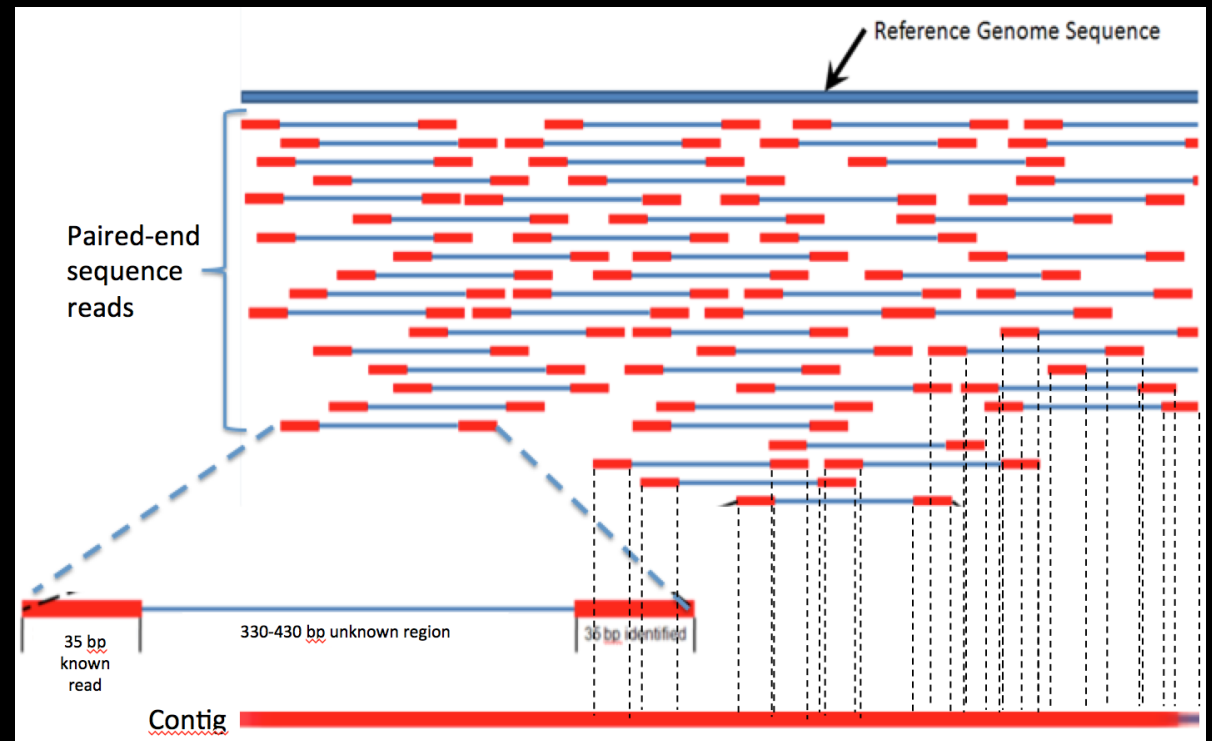


Paired-End



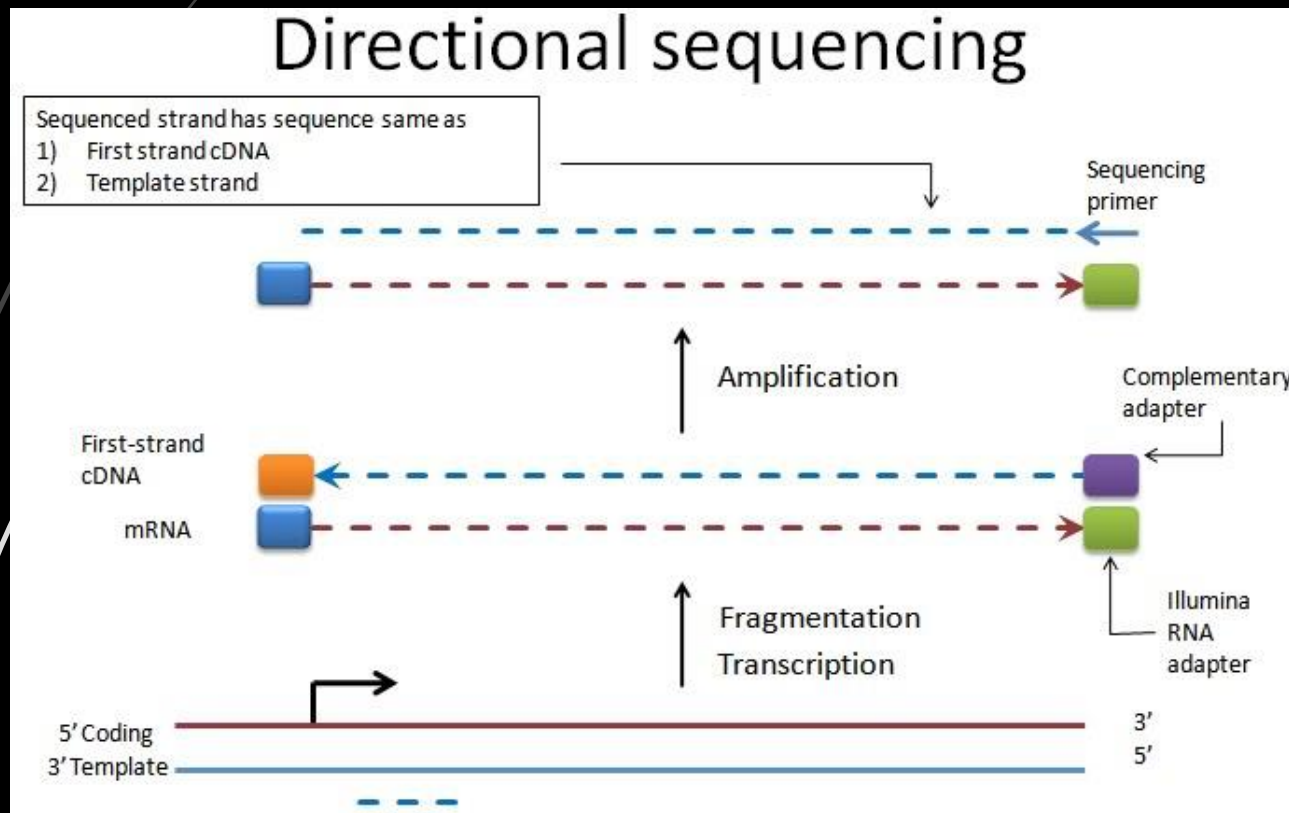
Fragment size = R1+R2+insert size

PE reads: Higher alignment rate and accuracy



2 x 50bp PE reads >> 1 x 100 bp SE reads!

Library strandness: stranded-protocol is always recommended!




- Non-stranded/Non-directional
- Forward, also known as second-stranded (coding strand)
- Reverse, also known as first-stranded (template strand)

Sequencing depth: How many reads do I need?

Coverage affects downstream gene discovery, expression estimation, and power of statistical analysis

The more, the better!



Trapnell et al. Nature Protocol 2012

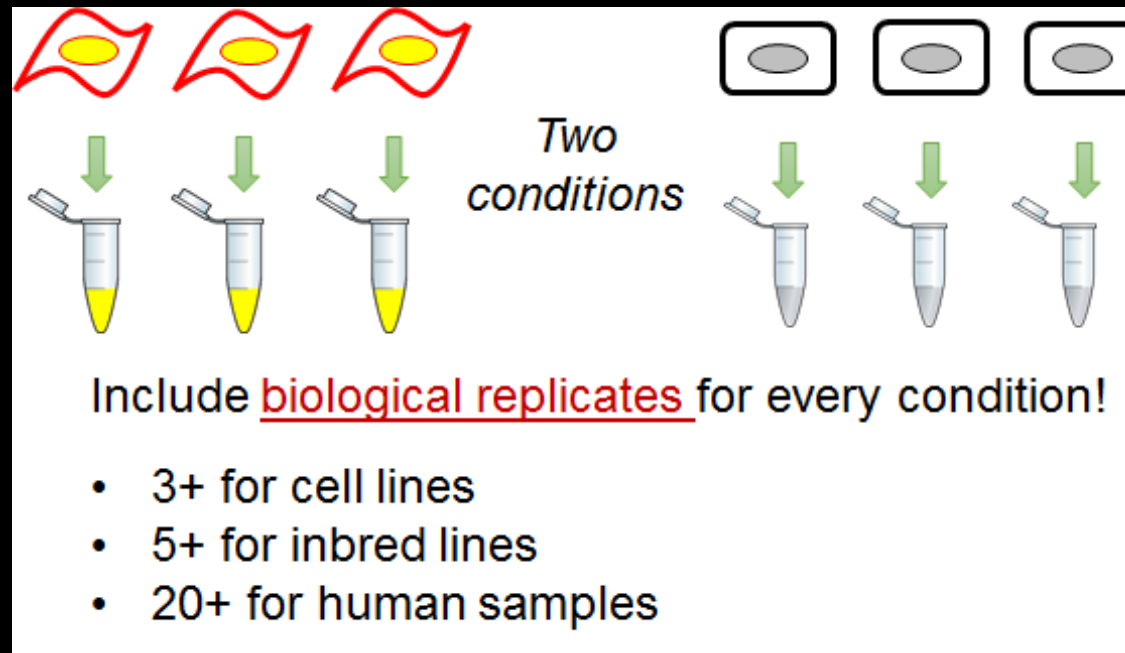
ENCODE saturation analysis

- 214 million 2x100bp PE reads
- H1 human embryonic stem cells
- 80% of the genes with $\text{FPKM} \geq 10$ are detected by ~36 million mapped reads per sample
- Genes with $\text{FPKM} < 10$: ~80 million mapped reads per sample

Sims et al., 2014 Nature Reviews


Experimental design: Biological replicates

Include biological replicates for **increased discovery power** and reduced false positives/negatives!



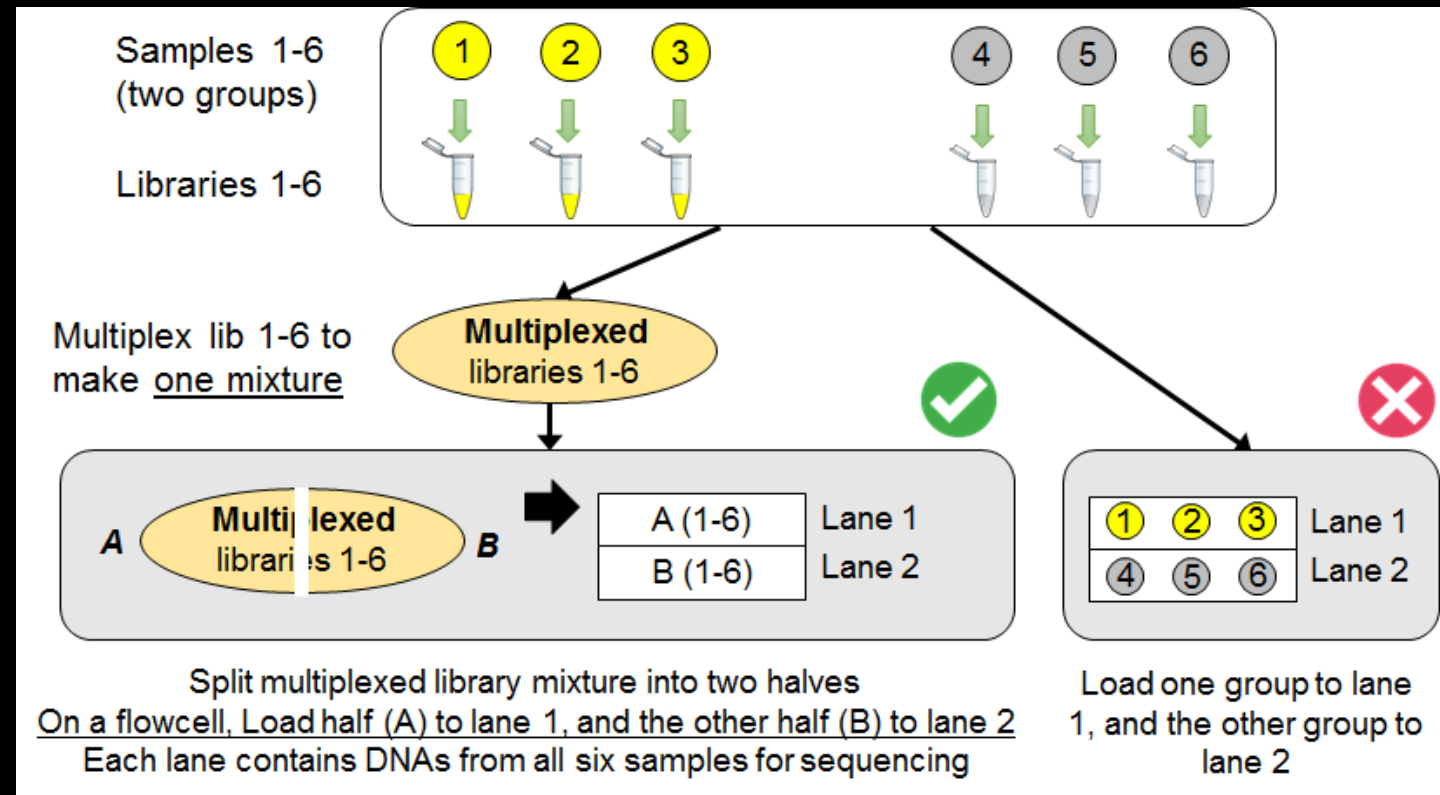
Experimental design: Multiplexing and Randomization

- **Multiplexing:** simultaneously measures multiple libraries in one sequencing lane. Unique barcodes are added to label DNA molecules from each library
- **Randomization:** Avoid loading samples from the same biological group in the same sequencing lane. Minimizes technical bias and lane-specific effects.




Sequencing lanes


Experimental design: Multiplexing and Randomization



Challenges and limitations

- Relatively poor RNA quality for tumor FFPE samples
- Contamination from adjacent normal tissue
- Still more expensive than targeted-panel sequencing such as NanoString
- 40 million mapped reads are usually sufficient for gene profiling, but > 80 million are required to detect bottom 1% lowly expressed genes





Whom to contact if I want to initiate an NGS (e.g. RNAseq) experiment?

- The University of Chicago Genomics Facility

Pieter W. Faber, Ph.D.

Technical Director

pfaber@bsd.uchicago.edu

773.834.8420

- The Center for Research Informatics Bioinformatics Core

Jorge Andrade, Ph.D.

Director of Bioinformatics




jandrade@bsd.uchicago.edu

How to perform RNAseq analysis

The good-practice analysis protocol takes 8 major steps.

- **01-04:** From raw sequencing to transcript quantification
- **05-08:** DEG and pathway analysis (06/03, part II)

```
CTTTGGATGACTTCACA  
CTAGCTTCCGCTTTCTT  
CTGCTCAAAAGTCTTCA  
TCTTGATTTGACCAGTT  
CTTAACACAGTTGCATA  
ATTTAGGTTACAATTTA  
TTGCTTGATATCCACCA
```

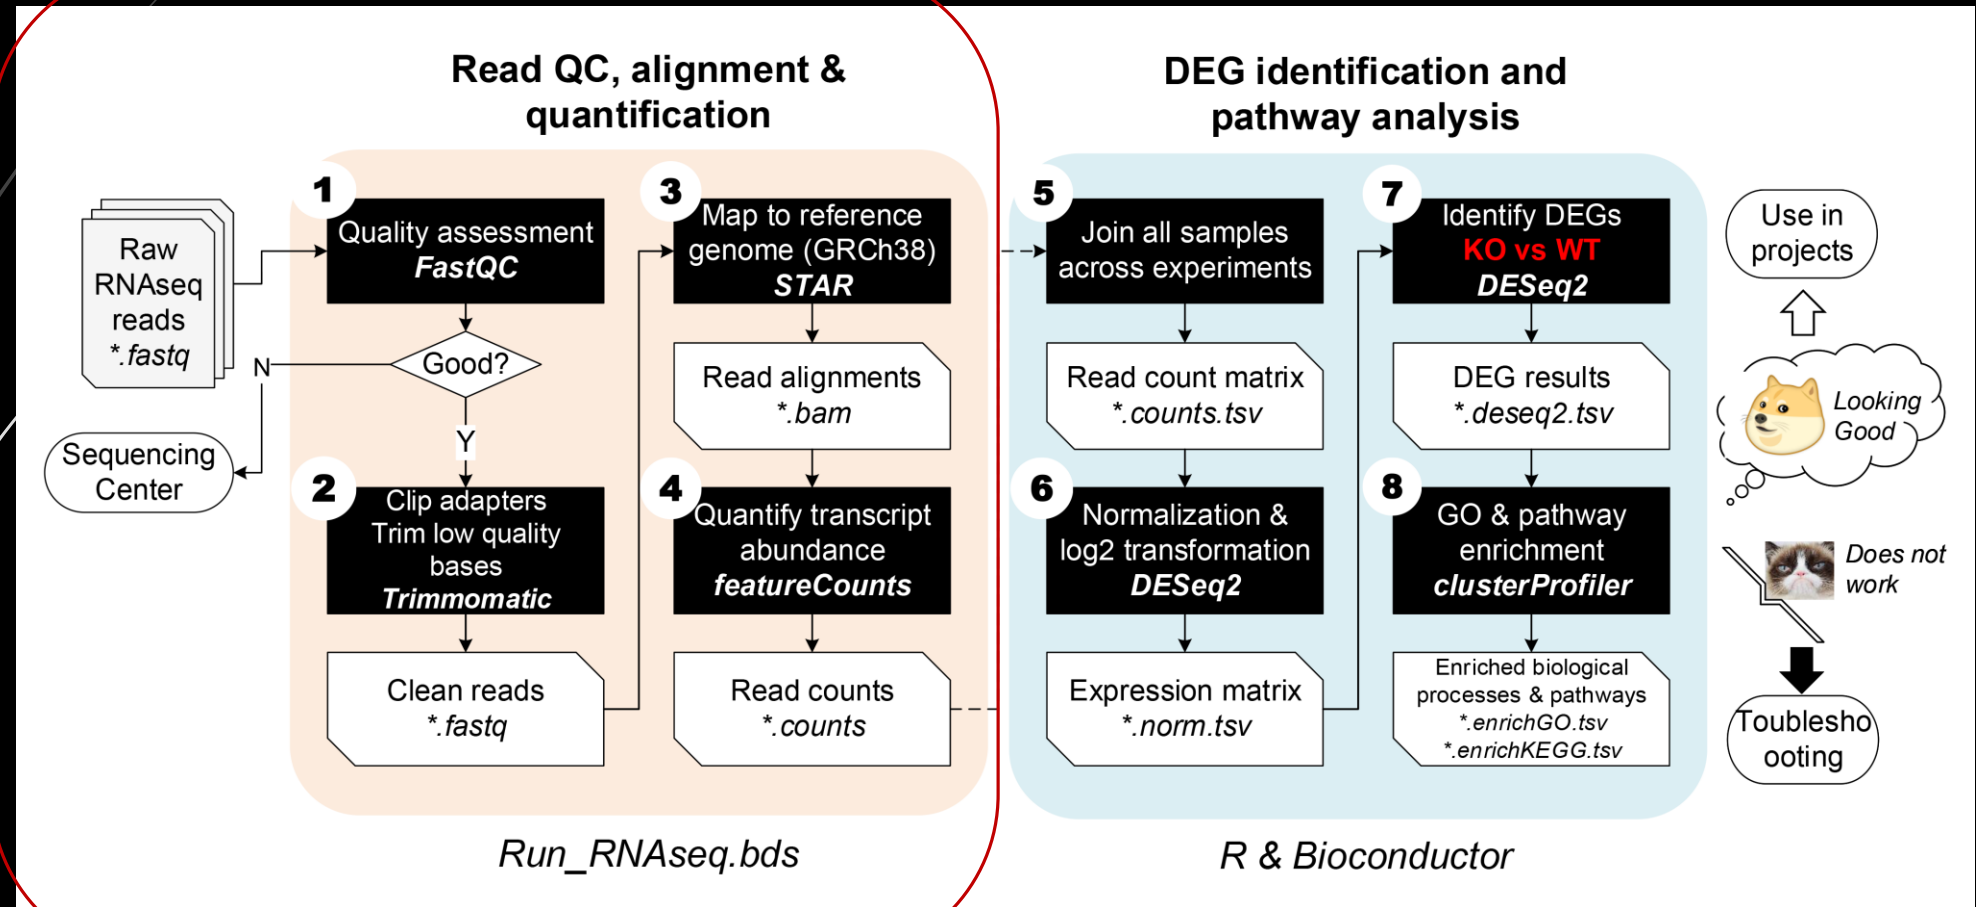


Raw sequencing data + sample group



Differentially expressed genes and pathways

How to perform RNAseq analysis



01 Quality assessment of raw sequencing reads: FastQC

- Raw sequencing reads are stored in **FastQ** format (e.g. *KO01.fastq.gz*), where each read is presented by 4 lines

Sequence header --	@SRR1205282.43583628
Sequence string --	GACTATCTTGGCCAACATGGTGAAACCCCGTCTCTACTAAAAATACAAA
Quality header --	+
Quality string --	CCCFFFFFHHHHHHIIIIIIHHIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII

- Check if the reads are of high quality
- Check if any preprocessing step is required (e.g. base trimming, adapter clipping, read filtering)



01 Quality assessment of raw sequencing reads: FastQC

- ▶ Method
 - ▶ FastQC version 0.11.5
 - ▶ Scan raw sequencing reads and produce QC reports for evaluation
- ▶ Our data
 - ▶ Read quality pretty good (baseQ \geq 30 in all base positions)
 - ▶ Preprocessing is optional

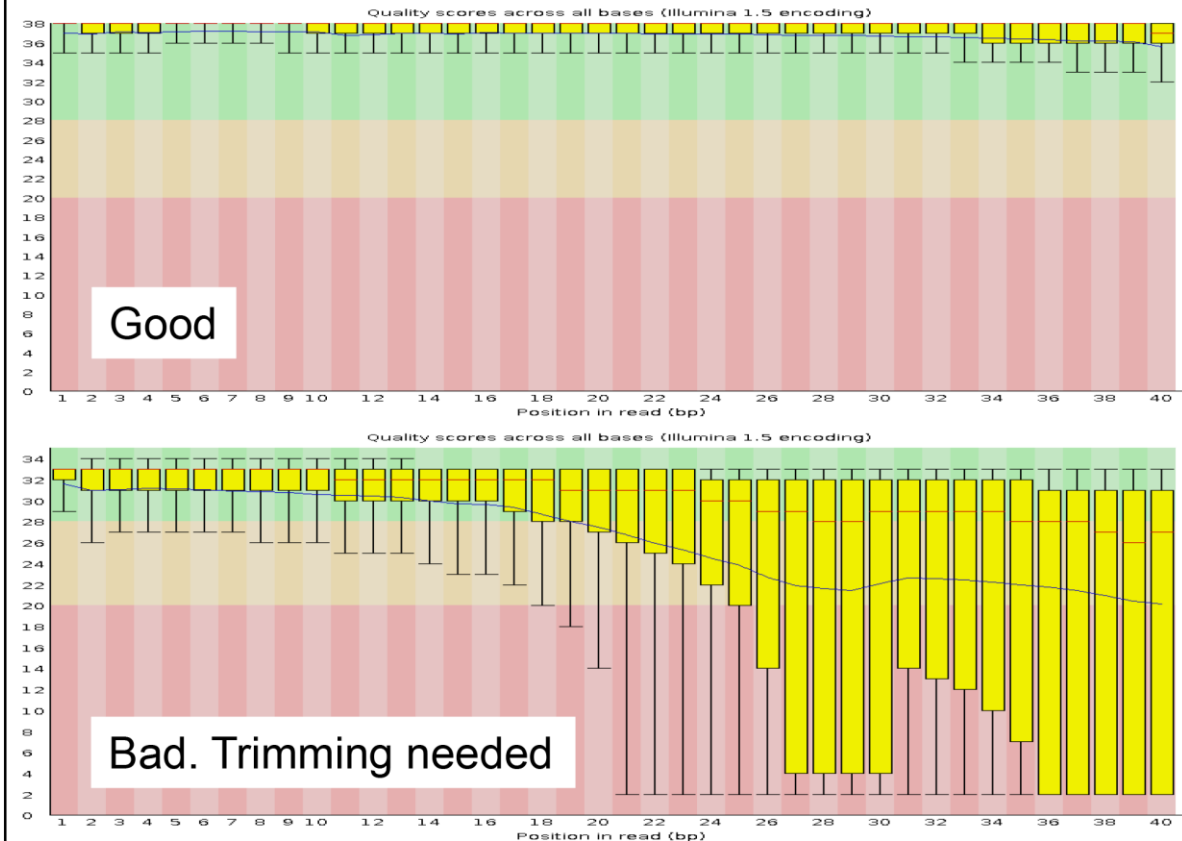
```
fastqc --extract -o $out.dir -t 2 --nogroup $r1.fastq.gz
```

01 Quality assessment of raw sequencing reads: FastQC

Summary

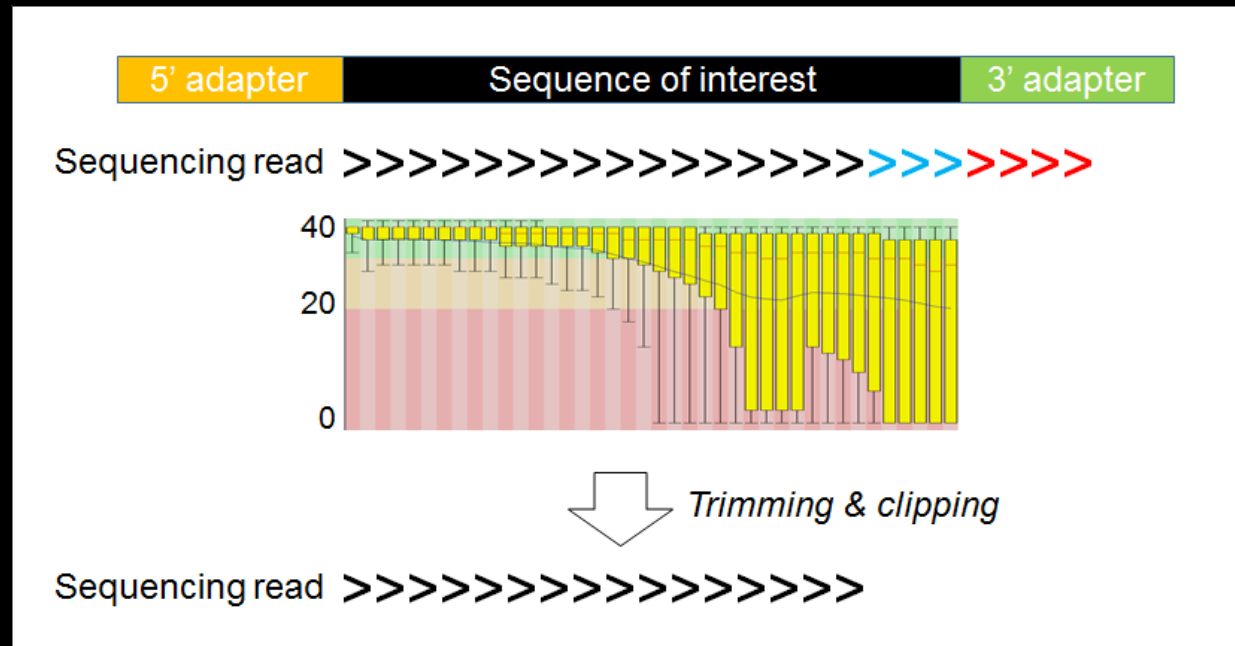
- ✓ [Basic Statistics](#)
- ✓ [Per base sequence quality](#)
- ! [Per tile sequence quality](#)
- ✓ [Per sequence quality scores](#)
- ✗ [Per base sequence content](#)
- ✓ [Per sequence GC content](#)
- ✓ [Per base N content](#)
- ✓ [Sequence Length Distribution](#)
- ✗ [Sequence Duplication Levels](#)
- ✓ [Overrepresented sequences](#)
- ✓ [Adapter Content](#)
- ✗ [Kmer Content](#)

Per base sequence quality



02 Preprocessing: Trimmomatic

- Preprocess reads to improve mapping rate and accuracy
 - Trim low-quality bases, clip adapters, etc.
 - Avoid over-trimming in RNAseq!
- Clean up reads for improved alignment rate and accuracy (for the next step)



02 Preprocessing: Trimmomatic

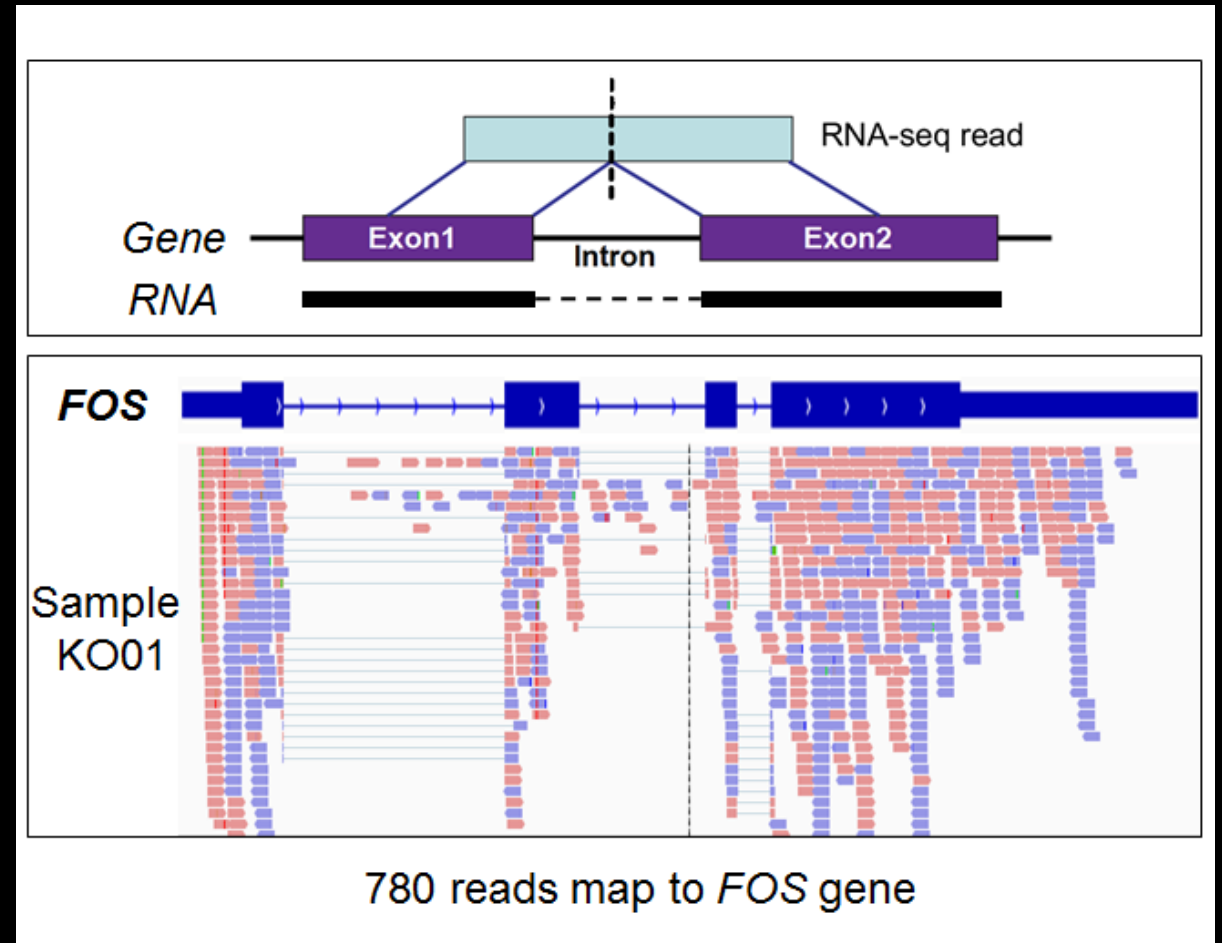
- ▶ Method
 - ▶ Trimmomatic version 0.36
 - ▶ Clip adapters
 - ▶ Trim leading/trailing low quality or N bases
 - ▶ Trim reads to a specific length
 - ▶ Filter out reads of low average quality / of specific length
 - ▶ Convert base quality scores between Phred33 and Phred64 FastQ format
- ▶ Our data (KO01 as an example)
 - ▶ 74,126,025 reads total. Survived: 73,636,793 (99.34%) Dropped: 489,232 (0.66%)

```
java -Xmx4G -jar trimmomatic-0.36.jar SE -threads 4 -phred33  
$r1.fastq.gz $r1.trim.fastq.gz ILLUMINACLIP:TruSeq3-  
SE.fa:2:30:10 LEADING:5 TRAILING:5 MINLEN:36 SLIDINGWINDOW:4:15
```

03 Map reads to reference genome (GRCh38): STAR

- Read mapping identifies the location in the genome where a sequencing read comes from
- **Splice-aware aligner** (e.g. *STAR*)

Each horizontal bar represents one read. Red/blue indicates reads aligned to plus/minus strand on the genome, respectively.



03 Map reads to reference genome (GRCh38): STAR

- ▶ A read may map to one, or multiple genomic locations (Mapping quality: MAPQ)
- ▶ Accurate mapping result is the key for downstream differential gene expression identification
- ▶ Read alignment results are stored in **SAM** or **BAM (Binary SAM)** format

HWI-ST330:304:H045HADXX:20934#0	16	chr1	60023	50	100M	*	0	0
CCACTATGTTTTTCGATAAAAAGCTTAATAAAT	?????BBBBBDBDB=?FFECFACCFHHH>09C							
↑	↑	↑	↑	↑	↑			
QNAME	FLAG	RNAME	POS	MAPQ	CIGAR			
HWI-ST330:304:H045HADXX:20934#0	16	chr1	60023	50	100M			

HWI-ST330:304:H045HADXX:20934#0	16	chr1	60023	50	100M	*	0	0
CCACTATGTTTTTCGATAAAAAGCTTAATAAAT	?????BBBBBDBDB=?FFECFACCFHHH>09C							
↑	↑	↑	↑	↑	↑	↑	↑	↑
SEQ	QUAL	MRNM	MPOS	ISIZE				
		*	0	0				

03 Map reads to reference genome (GRCh38): STAR

- Method
 - STAR version 2.5.3a
 - splice-aware aligner
 - Ultrafast (~15 minutes for 50m reads), require lots of memory (~36GB for human genome)
 - Flexible options to allow canonical/non-canonical junctions, with/wo known gene annotations, etc.
- Different aligners may generate very different results! (Engström et al. *Nature methods* 2013. Systematic evaluation of spliced alignment programs for RNA-seq data)

```
STAR --runMode alignReads --genomeLoad NoSharedMemory --  
outFileNamePrefix $out.prefix --readFilesCommand zcat --  
genomeDir $refgenome.dir --readFilesIn $r1.trim.fastq.gz --  
runThreadN 2 --outSAMstrandField intronMotif --  
outFilterIntronMotifs RemoveNoncanonicalUnannotated --  
outSAMtype BAM SortedByCoordinate
```

Reference
genome

KO01
coverage

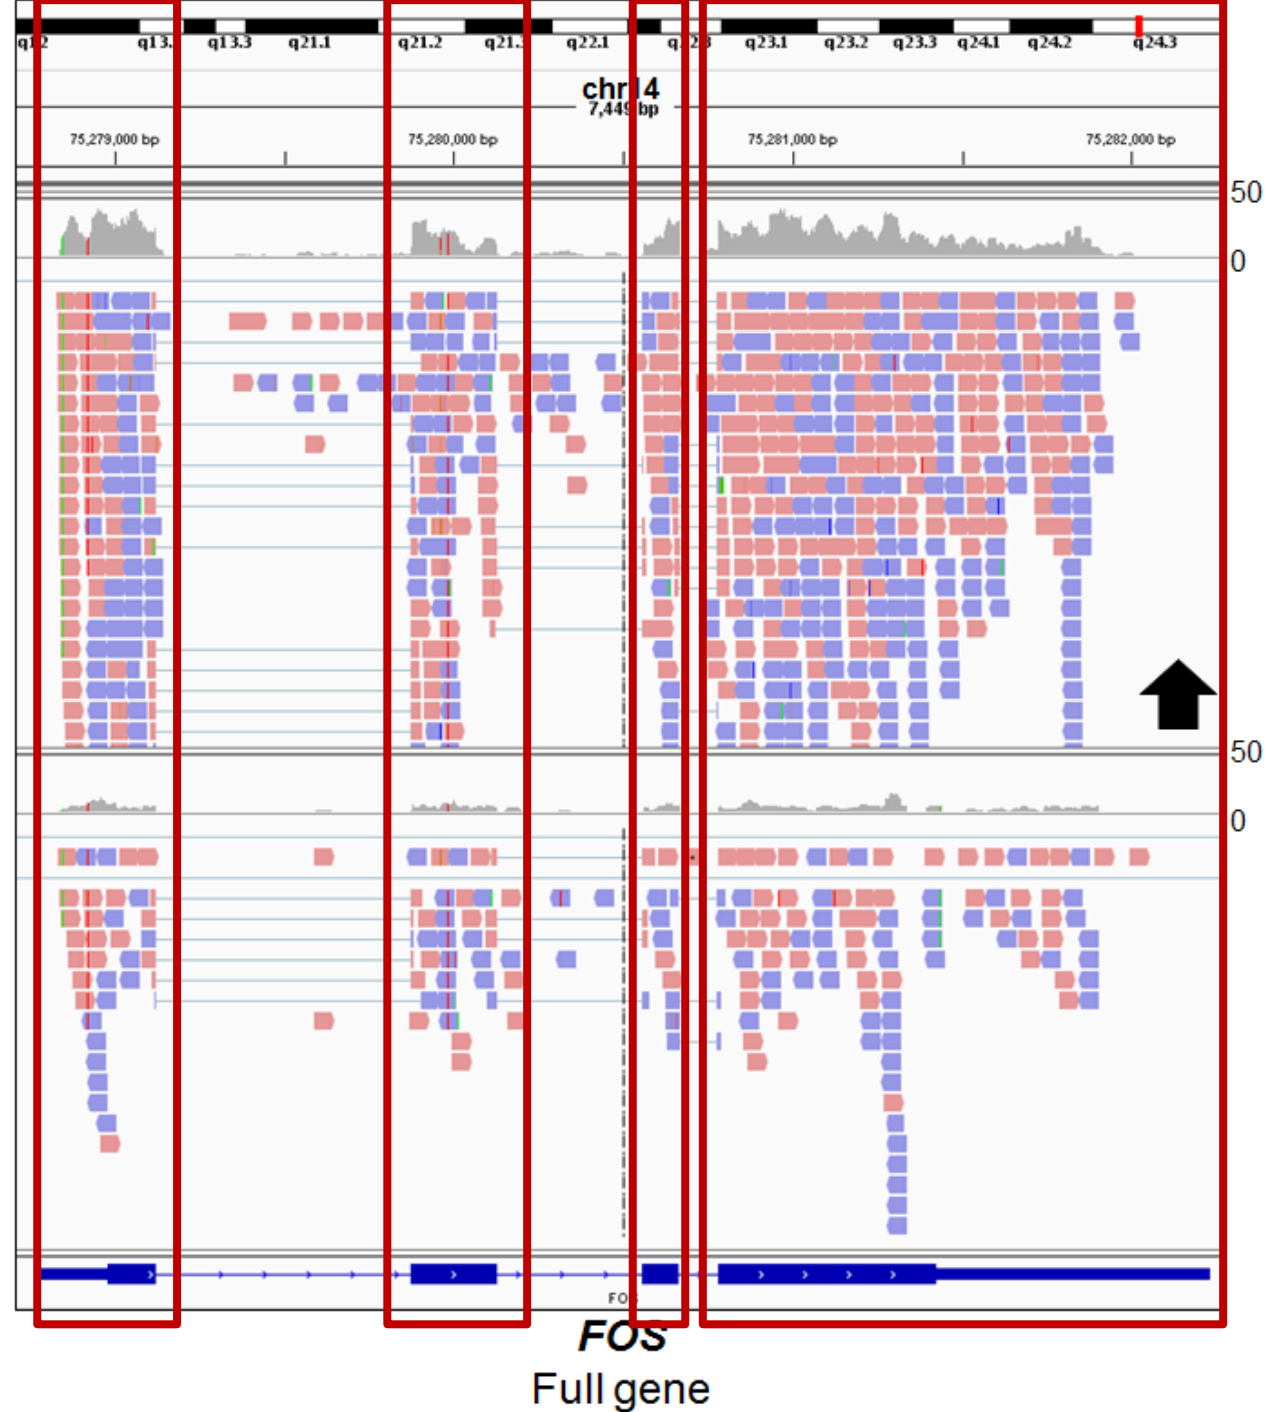
KO01
alignment

siPRDM11
knockdown

WT01
coverage

WT01
alignment

Gene



04 Quantify transcript abundance: featureCounts

- Estimate number of reads mapped to gene features (e.g. gene, exon, etc.)
- Method
 - featureCounts version 2.5.2b
 - Ultrafast (~10 minutes for 50m reads), require low amount of memory (~4GB for human genome)
 - Flexible options to count the reads based on specific mapping criteria or study purposes
 - Gene-level or exon-level (for isoforms); Uniquely mapped reads; Primary alignment; Properly paired reads (if reads are paired-end); Mapping quality thresholds
- Choose the option accordingly based on your experimental design!

```
featureCounts -s 0 -F GTF -t exon -g gene_id -Q 255 -J --primary -a  
$refgeneanno.gtf -T 2 -G $refgenome.fa -o  
$sample.star.featurecounts.raw_counts.txt $sample.star.merged.bam
```

```
//===== featureCounts setting =====//
Input files : 1 BAM file
               S results/rnaseq/DLBC_samples_grch38/K001/al ...

Output file : results/rnaseq/DLBC_samples_grch38/K001/read ...
Annotations : gencode.v24.primary_assembly.annotation.chr11...
Junction Counting : <output_file>.jcounts

Threads : 12
Level : meta-feature level
Paired-end : yes
Strand specific : no
Multimapping reads : primary only
Multi-overlapping reads : not counted
Read orientations : fr

Chimeric reads : not counted
Both ends mapped : not required

//===== http://subread.sourceforge.net/ =====//

//===== Running =====//
Load annotation file gencode.v24.primary_assembly.annotation.chr11.gtf
Features : 1177678
Meta-features : 60725
Chromosomes/contigs : 59

Loading FASTA contigs : GRCh38.primary_assembly.genome.chr11.fa
194 contigs were loaded

Process BAM file results/rnaseq/DLBC_samples_grch38/K001/alignment/K00 ...
Single-end reads are included.
Assign reads to features...
Total reads : 93186698
Successfully assigned reads : 55910832 (60.0%)
Running time : 10.23 minutes

Found 230830 junctions in all the input files.


Read assignment finished.

//===== http://subread.sourceforge.net/ =====//
```


Geneid	Chr	Start	End	Strand	Length	KO01
ENSG00000019485.12	chr11;chr11;chr11;chr11;chr11;chr11;chr11;chr11;chr11;chr11	45095806;45146675;45147343;45181761;45182246;45182861;45204711;45208874;45212469;45219570;45224217;45225995	45095901;45146877;45147581;45181885;45182349;45183123;45204778;45209295;45214713;45219757;45225087;45235124	+++++; +++	13954	710
ENSG000000170345.9	chr14;chr14;chr14	75278774;75279237;75279643	75279128;75279531;75282230	+++	3238	780
ENSG000000119660.4	chr14	75292131	75292495	-	365	0
ENSG000000259687.1	chr14;chr14;chr14	75294404;75294677;75296120	75294541;75294737;75296638	+++	718	0
ENSG000000259319.1	chr14	75423683	75427741	-	4059	25


Pseudoaligners (Kallisto, Salmon)

- 'Align' reads to the transcriptome instead of genome
- K-mer 'compatibility' searching (pseudo-alignment) *BAM file optional*
- Accurate transcript-level quantification, robust to sequencing errors
- Gene-level quantification can be summarized by `tximport` (subsequent to `kallisto`) or `salmon` itself
- Super fast and relatively low memory requirement!



Transcriptome
de Bruijn graph
(T-DBG)



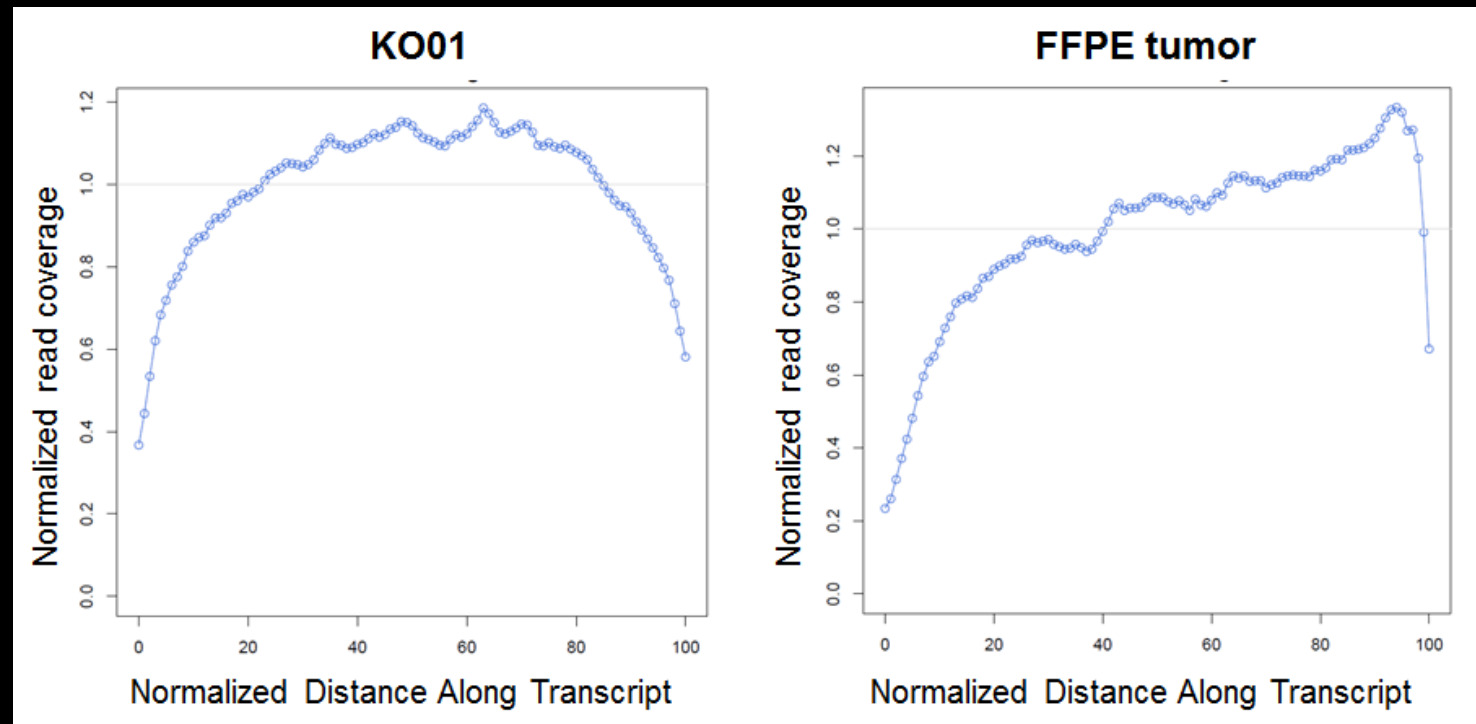


RNAseq quality metrics & coverage: Picardtools, bedtools, RSeQC

- Evaluate the quality of reads and alignments
- Identify potential problems regarding (RNA) sample quality
 - Is the RNA highly degraded?
 - Is there high-level genomic DNA contamination?
 - Was ribosome RNA successfully depleted?
 - How do reads distribute on the genome? (exons, introns, intergenic, etc.)
 - Is the strandness of read alignment consistent with library type? non-stranded or forward/reverse strand-specific libs
 - Does the target gene that was knocked down in KO samples have reduced expression, compared to WT?

a) RNA degradation

- ▶ RNA quality: Fresh samples (e.g. cell line) > frozen samples (e.g. mouse tissues) >> FFPE samples (e.g. human tumors)
- ▶ During lib prep, RNA quality is inferred by The RNA integrity number (RIN) evaluated using the 28S to 18S rRNA ratio (e.g. RIN > 4.5)
- ▶ However, studies have shown that RIN can be quite inaccurate for FFPE samples
- ▶ *Gene body coverage plot* (Picardtools). FFPE samples often has 5' degradation






b) Genomic DNA contamination

- ▶ In lib prep, DNase digestion removes genomic DNA
- ▶ While fresh samples often has good-quality RNAs, highly-degraded samples (e.g. FFPE tumors) often has a higher degree of genomic DNA contamination
- ▶ Sometimes DNA contamination could occupy 70% of sequencing reads, greatly reducing the discovery power of DEG analysis
- ▶ Good assessment to identify and estimate genomic DNA contamination includes
 - ▶ Read distribution in genomic features: high fraction of intergenic reads indicates DNA contamination
 - ▶ Visualize intergenic region in genome browser (e.g. IGV)

b) Genomic DNA contamination

Q1: Which sample (S1-4) has the most severe genomic DNA contamination?

Hint: higher percentage of intergenic reads indicates more severe DNA contamination in RNA samples






c) Ribosome RNA fraction

- ▶ rRNA accounts for > 80% of the transcriptome
- ▶ Most RNAseq lib prep protocol includes a ribosomal RNA depletion step
- ▶ However, if RNA quality is relatively poor (e.g. FFPE tumors), rRNA depletion often is efficient
- ▶ Assessing if the depletion step is successful through read distribution in genomic features
- ▶ How many reads map to ribosome RNA regions on the genome?


c) Ribosome RNA fraction

Q2: Which sample (S1-4) has the least efficient depletion of ribosome RNAs?

Hint: rRNAs account for > 80% of the whole transcriptome. If not removed, the majority of the sequencing reads will be derived from rRNA



c) Library standness: fraction of correctly oriented reads



RSeQC: infer_experiment.py

For pair-end RNA-seq, there are two different ways to strand reads (such as Illumina ScriptSeq protocol):

1. 1++,1-,2+-,2-+

- read1 mapped to '+' strand indicates parental gene on '+' strand
- read1 mapped to '-' strand indicates parental gene on '-' strand
- read2 mapped to '+' strand indicates parental gene on '-' strand
- read2 mapped to '-' strand indicates parental gene on '+' strand

2. 1+-,1-+,2++,2--

- read1 mapped to '+' strand indicates parental gene on '-' strand
- read1 mapped to '-' strand indicates parental gene on '+' strand
- read2 mapped to '+' strand indicates parental gene on '+' strand
- read2 mapped to '-' strand indicates parental gene on '-' strand

For single-end RNA-seq, there are also two different ways to strand reads:

1. ++,-

- read mapped to '+' strand indicates parental gene on '+' strand
- read mapped to '-' strand indicates parental gene on '-' strand


2. +,-+

- read mapped to '+' strand indicates parental gene on '-' strand
- read mapped to '-' strand indicates parental gene on '+' strand



e) Confirmation of reduced/increased expression for knockdown/overexpressed genes in KO/OE samples

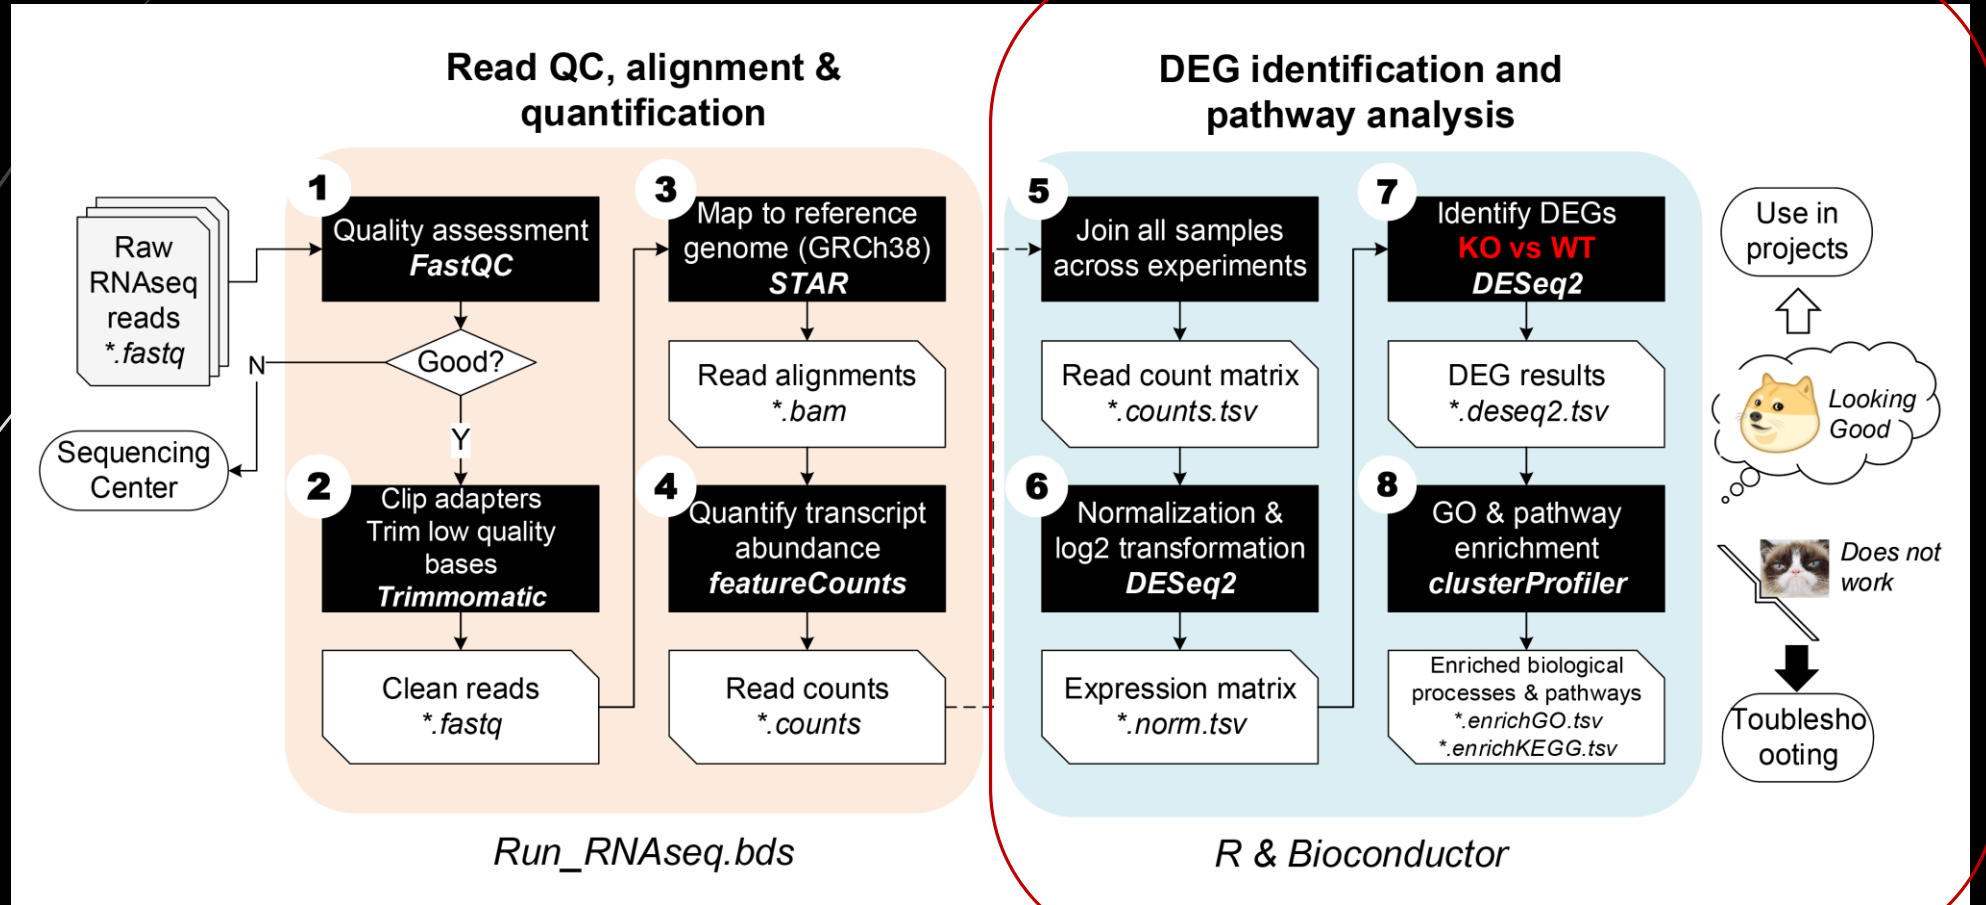
- Is there a gene that is expected/known to be repressed or overexpressed?
- Does the RNAseq result reflect expression change of this gene?
- In our data, *PRDM11* expression is expected to be down since this is the gene knocked down in U2932 cells
- NOT all exons are affected!
- Only those targeted by siRNA (small RNA interference) will be affected



➤ Our data


- *PRDM11* knockdown U2932 cells in triplicates (KO01-03 vs WT01-03)
- NOT the full *PRDM11* gene is knockdown!
- siRNAs target exons 4,5,6 of *PRDM11*, thus only those three exons show a reduction of expression in the KO samples; other exons are not affected
- *PRDM11* knockdown leads to upregulation of *FOS* expression

How to perform RNAseq analysis





05-08: Identify differentially expressed genes and pathways: DESeq2, clusterProfiler



- After steps 01-04, we have generated read alignment and counts for every annotated gene on the genome
- The next step is to utilize the read counts data to detect DEGs
- For example, if we visualize *FOS* gene across 6 samples in genome browser




FOS = *Fos* proto-oncogene, AP-1 transcription factor subunit

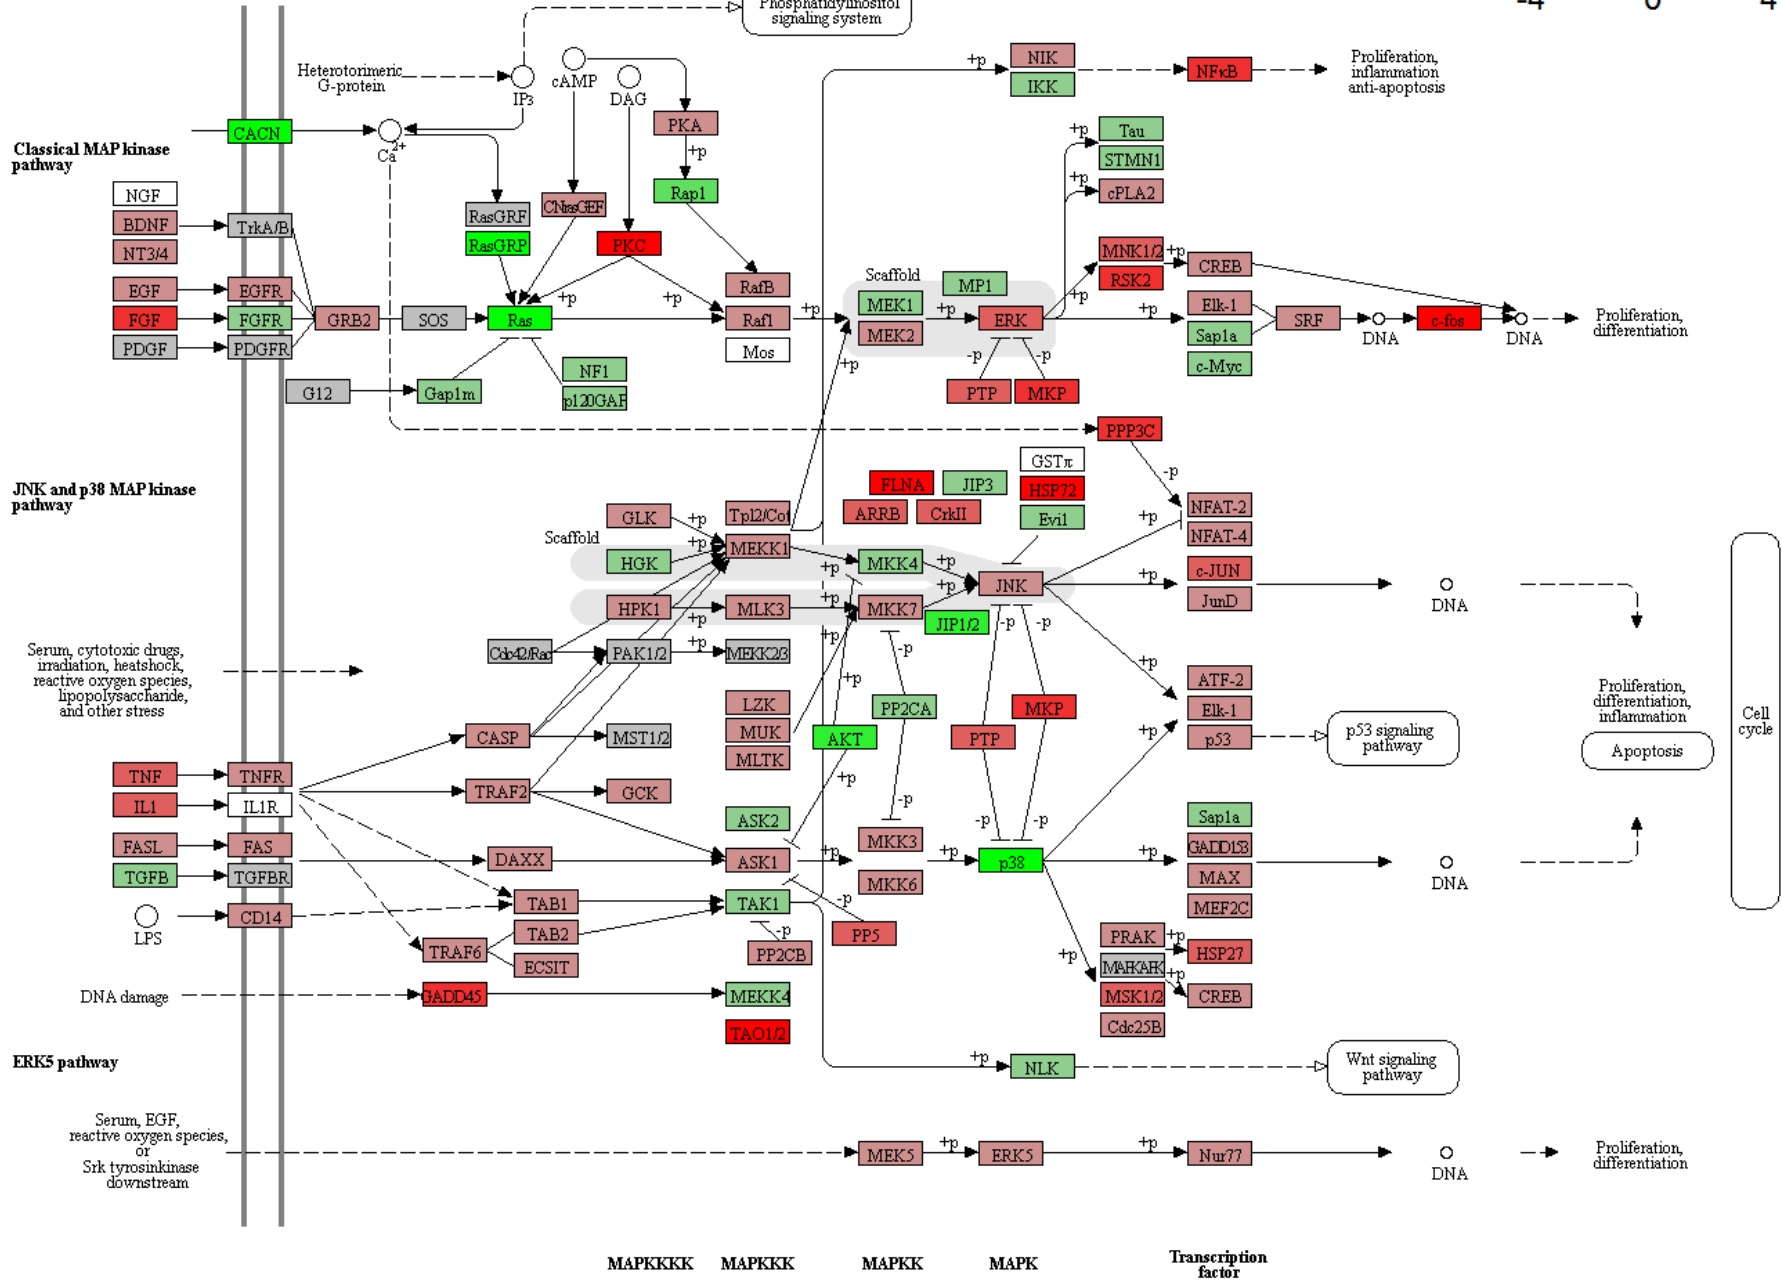



[1] "Genes significant = 296 (fc, 1.5, fdr 0.05)"
 [1] "Heatmap = 296 genes on the row, 6 samples on the column"






MAPK SIGNALING PATHWAY



Thank you!



Questions



Hands-on practice START


- Open your handson.Rmd on the Github or download to local computer
- <https://github.com/MScBiomedicalInformatics/MSIB32500/blob/master/lectures/handson8.html>
- Dataset: two groups (PRDM11 KO vs WT, human U2932 cells), 6 samples
- Single-end reads, unstranded libraries

Sample	Group	Sequencing File	Sequencing Data
KO01	KO	KO01.fastq.gz	74,126,025 reads
KO02	KO	KO02.fastq.gz	64,695,948 reads
KO03	KO	KO03.fastq.gz	52,972,573 reads
WT01	WT	WT01.fastq.gz	55,005,729 reads
WT01	WT	WT02.fastq.gz	61,079,377 reads
WT01	WT	WT03.fastq.gz	66,517,156 reads

Fog. et al. 2015. Loss of *PRDM11* promotes MYC-driven lymphomagenesis. Blood 125(8):1272-81

PRDM11 = PR/SET domain 11

Job running is prohibited on the head node!





Ten essential Linux commands

cd	Change directory
ls	List contents
cp	Copy
mv	Move/rename
rm	Delete
pwd	Print the current path
head	Show the first few lines of a file
more	View a file by page
mkdir	Create a new directory
man	Show help menu for a command



Five useful HPC/Shell commands

qsub	Submit a job
qstat	List submitted jobs from the user
qdel	Delete a job
showq	Show all jobs on the cluster
ssh	Log into a server using Secure Shell (SSH)

